

BIL216 — Signals and Systems

Final Project Phase 2 Report

Emo Challenge 2026 — Speech Emotion Classification

Course	BIL216 — Signals and Systems
Semester	2025–2026 Spring
Group	9
Submission	Phase 2 — Research & Development

Group Members

Name	Student ID	Role
Ahmet Akin	230611038	Feature Engineering & Software Implementation
Berivan Demir	230611030	Model Design, Ensemble & Hyperparameter Optimisation
Mustafa Talha Akgul	230611043	Error Analysis, Report & Leaderboard Submission

Submission Date: May 19, 2026

1. Introduction

In Phase 1 we built a baseline speech emotion classifier using a 51-dimensional handcrafted feature set and a single RandomForestClassifier. The system achieved a test accuracy of **62.6%** (macro-F1 = 0.626) on the five-class problem (*Neutral, Happy, Angry, Sad, Surprised*). Error analysis revealed three weaknesses: (i) the model could not separate high-arousal emotions (*angry* and *happy* were frequently confused), (ii) *surprised* was systematically under-detected (recall 0.333) because its short pitch-rise dynamics were not captured by frame-aggregate statistics, and (iii) low-arousal classes (*sad* and *neutral*) overlapped in the energy space.

Phase 2 — **Research & Development** — addresses these weaknesses with three concrete upgrades motivated by the speech-emotion-recognition (SER) literature:

1. **Feature expansion** from 51 → **135 dimensions**, adding temporal-derivative MFCCs (Δ -MFCC), tonal features (Chroma, Tonnetz), Mel-spectrogram statistics, and spectral flatness.
2. **Multi-model comparison**: alongside an upgraded RandomForest we train **XGBoost** and a **multi-layer perceptron (MLP)**, then fuse them with a **soft-voting ensemble**.
3. **Hyper-parameter optimisation** via **RandomizedSearchCV** on each base learner.

The Phase 2 ensemble reaches **89.3% test accuracy** (macro-F1 = 0.891) — an absolute improvement of **+26.7 percentage points** over the Phase 1 baseline. The most striking gain is on *surprised*, whose recall climbs from **0.333** → **0.778**.

2. Literature Review

The features and methodology of Phase 2 were chosen after reviewing the following references:

1. **Schuller, B. et al.** *"Speech Emotion Recognition: Two Decades in a Nutshell, Benchmarks, and Ongoing Trends,"* Comm. ACM, 2018. Establishes the gold-standard handcrafted feature taxonomy (prosodic, spectral, voice-quality) and reports that ensemble classifiers consistently outperform single learners on the IEMOCAP and EmoDB benchmarks.
2. **Furui, S.** *"Speaker-Independent Isolated Word Recognition Based on Emphasized Spectral Dynamics,"* ICASSP 1986. Original motivation for **delta cepstra**: first-order time derivatives of MFCC frames encode rate-of-change information that the static cepstral coefficients alone discard.
3. **Bou-Ghazale, S. E. and Hansen, J. H. L.** *"A Comparative Study of Traditional and Newly Proposed Features for Recognition of Speech Under Stress,"* IEEE Trans. Speech Audio Process., 2000. Demonstrates that Δ -MFCC and Δ^2 -MFCC together yield 10–15% relative improvement on stressed/emotional speech versus static MFCC alone.
4. **Eyben, F. et al.** *"The Geneva Minimalistic Acoustic Parameter Set (GeMAPS) for Voice Research and Affective Computing,"* IEEE Trans. Affect. Comput., 2016. Defines a compact, well-justified feature set for emotion that prominently includes **spectral flatness**, **MFCC dynamics**, and **harmonic features** — all of which we include.
5. **Livingstone, S. R. and Russo, F. A.** *"The Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS),"* PLOS ONE, 2018. Reference dataset whose published baselines (MLP, RF) we used as a sanity check for our accuracy range.

6. **Wang, K. et al.** "Speech Emotion Recognition Using Fourier Parameters and Random Forest Ensembles," *Multimed. Tools Appl.*, 2015. Empirical evidence that **soft-voting ensembles** of heterogeneous classifiers (tree-based + neural) improve recall on minority classes — exactly the problem we observed for *surprised* in Phase 1.

The combined message of these works informed three design choices: (a) extend the feature set with temporal dynamics and tonal descriptors, (b) replace the single RF with an ensemble of heterogeneous learners, and (c) use cross-validated hyper-parameter search to fairly compare them.

3. Dataset Characterization

The dataset is unchanged from Phase 1: ~675 WAV recordings collected by all class groups during the Midterm Project, labelled into five emotions. Files follow the naming convention

G<GroupID>_D<SpeakerID>_<Gender>_<Age>_<Emotion>_C<Quality>.wav

Statistic	Value
Total recording groups	39
Total valid WAV files	~675
Emotion classes	5 (Neutral, Happy, Angry, Sad, Surprised)
Approximate files per class	~135
Age range	8 – 87 years
Speaker types	Male / Female / Child
Language	Turkish (native speakers)
Sample rate (after resample)	22,050 Hz

Turkish emotion labels (e.g., *Notr*, *Mutlu*, *Öfkeli*, *Üzgün*, *Şaşkın*) are normalised to the five canonical English classes at load time via a lookup table. Group `GRUP_20` was excluded because it contained no usable WAV files (only `.m4a.mp4` and `.ogg`).

Train/test split is held identical to Phase 1 (80/20, stratified, `random_state=42`) to make the comparison fair: any accuracy gain is attributable to features or models, not to a different split.

4. Methodology — Phase 2 Enhancements

4.1 Expanded Feature Set (51 → 135 dims)

Phase 1 features are retained in full. Five new feature groups (84 additional dimensions) are appended.

Feature Group	Dims	Captures	Motivation
MFCC mean+std (13 coeff.)	26	Static spectral envelope (vocal tract shape)	Phase 1 — kept
STE mean, ZCR mean+std	3	Loudness, noisiness	Phase 1 — kept
F0 mean+std (autocorrelation)	2	Pitch / prosody	Phase 1 — kept
Spectral Centroid/BW/Rolloff (mean+std)	6	Brightness, spread, energy concentration	Phase 1 — kept
Spectral Contrast (7 bands, mean+std)	14	Peak–valley contrast across sub-bands	Phase 1 — kept
Δ-MFCC mean+std (13 coeff.)	26	First-order temporal derivative of cepstra	NEW: encodes <i>rate of change</i> of vocal tract — captures the rapid pitch-rise of <i>surprised</i> and the irregular voice quality of <i>anger</i> (Furui 1986).
Chroma STFT (12 pitch classes, mean+std)	24	Tonal / harmonic energy distribution	NEW: emotion correlates with harmonic content — <i>happy</i> tends to major intervals, <i>sad</i> to minor (Eyben GeMAPS).
Spectral Flatness (mean+std)	2	Tonal-vs-noise ratio (Wiener entropy)	NEW: <i>angry</i> speech has low flatness (tonal); whispered or

Feature Group	Dims	Captures	Motivation
			<i>neutral</i> speech approaches white noise.
Mel-Spectrogram band statistics (10 bands × mean+std)	20	Coarse perceptual frequency envelope	NEW: Mel scale matches human auditory perception; pooled stats capture gross spectral shape that complements DCT-decorrelated MFCC.
Tonnetz (6 dims × mean+std)	12	Tonal centroid — projections onto perfect-fifth, minor-third, major-third axes	NEW: captures harmonic tension; <i>surprised</i> exhibits sudden tonal shifts that diverge from a neutral baseline.
TOTAL	135		

All 135 features are computed per-file via `librosa` native functions, then aggregated (mean + std across frames) to yield a fixed-length vector regardless of recording duration. The full preprocessing chain (preemphasis with $\alpha=0.97$ followed by `librosa.effects.trim(top_db=20)`) is unchanged from Phase 1.

4.2 Three Heterogeneous Classifiers

We deliberately choose three learners with **different inductive biases** so that their errors are partially decorrelated — a precondition for a useful soft-voting ensemble.

Classifier	Family	Key Hyper-parameters	Why this learner?
RandomForest	Bagged decision trees	<code>n_estimators=300</code> , <code>max_depth=20</code> , <code>max_features='sqrt'</code> , <code>class_weight='balanced'</code>	Robust to outliers; gives interpretable feature importance; serves as the Phase 1 anchor.
XGBoost	Gradient-boosted trees	<code>n_estimators=300</code> , <code>max_depth=6</code> , <code>learning_rate=0.1</code> , <code>subsample=0.8</code>	Sequential boosting corrects prior errors; built-in L1/L2 regularisation

Classifier	Family	Key Hyper-parameters	Why this learner?
			handles 135-dim feature collinearity better than vanilla RF.
MLP	Feed-forward neural network	hidden_layer_sizes=(256,128,64) , activation='relu' , solver='adam' , alpha=1e-3	Learns non-linear feature interactions (e.g., <i>high pitch</i> $\wedge$ <i>high energy</i> $\Rightarrow$ <i>anger</i>) that axis-aligned tree splits approximate only crudely.

4.3 Soft-Voting Ensemble

Given the three calibrated classifiers $f_{\text{RF}}, f_{\text{XGB}}, f_{\text{MLP}}$ each producing class-probability vectors $\mathbf{p}_i \in \Delta^4$ (five emotions), the ensemble prediction is

$$\hat{y}(x) = \arg \max_c \frac{1}{3} \sum_{i \in \{RF, XGB, MLP\}} p_i(c \mid x).$$

This averages probabilities rather than majority-voting hard labels, which is known to be strictly better when the base learners are well calibrated.

4.4 Hyper-Parameter Optimisation — RandomizedSearchCV

For each base learner we run an independent **3-fold cross-validated RandomizedSearchCV** on the training set with the following search spaces:

Model	Search-space size	Iterations sampled
RandomForest	$4 \times 4 \times 4 \times 3 \times 3 = 576$	20
XGBoost	$4 \times 5 \times 4 \times 4 \times 4 \times 4 \times 3 = 15,360$	20
MLP	$6 \times 4 \times 4 = 96$	15


RandomizedSearch with 20 samples gives ~95% probability of landing in the top 5% of the space — much cheaper than exhaustive GridSearch and empirically just as good for this problem size.

4.5 Validation

- 80/20 stratified train/test split (random_state=42 , identical to Phase 1).
- 5-fold StratifiedKFold cross-validation on the full scaled feature matrix for the ensemble.
- StandardScaler (zero-mean, unit-variance) is fit on the training portion only.

5. Statistical Findings

5.1 Feature Importance



The top of the ranking is dominated by **Δ -MFCC coefficients** (dMFCC_2_mean is the single most important feature, followed by dMFCC_2_std, dMFCC_5_mean, dMFCC_4_mean, dMFCC_7_mean, dMFCC_3_mean). This directly validates the Phase 2 hypothesis that *temporal dynamics* of the spectral envelope — absent in Phase 1 — carry decisive information for emotion. **Mel-spectrogram bands** also appear prominently (Mel_band9_std, Mel_band4_std, Mel_band2_std, Mel_band3_mean), confirming the value of perceptually-weighted frequency descriptors. Static MFCC means/stds (the Phase 1 work-horses) now appear lower in the ranking, indicating that the new derivative and Mel-spectrogram features explain variance the old features could not.


5.2 Acoustic-Profile-by-Emotion Validation

The expected acoustic profiles in Table 4.2 of the Phase 1 report (high F0 + high STE for *angry/happy*, low F0 + low STE for *sad*) are reflected in the classifier's behaviour: *sad* and *angry* are the two best-

recalled classes (0.941 and 0.909 respectively) while *neutral* — which sits in the middle of every prosodic axis — remains the most-confused class (recall 0.857, ~4/35 mis-classified as *happy*).

6. Classification Results

6.1 Phase 1 vs. Phase 2 — Model Comparison



Model	Feature dims	Test Accuracy	Δ vs. Phase 1
Phase 1 — RandomForest	51	62.6%	(baseline)
Phase 2 — RandomForest (tuned)	135	81.7%	+19.1 pp
Phase 2 — MLP (tuned)	135	89.3%	+26.7 pp
Phase 2 — Soft-Voting Ensemble	135	89.3%	+26.7 pp

Upgrading the feature set alone (RF: 51-dim → 135-dim) already accounts for ~19 pp of the gain. Switching to a stronger learner (MLP) extracts an additional ~8 pp. The ensemble matches MLP at


89.3% on the test set but is preferred because it is also more stable in cross-validation (lower fold-to-fold variance — see §6.5).

6.2 Per-Class Results (Soft-Voting Ensemble)

Class	Precision	Recall	F1-Score	Support
Angry	0.968	0.909	0.937	33
Happy	0.750	0.964	0.844	28
Neutral	0.968	0.857	0.909	35
Sad	0.889	0.941	0.914	17
Surprised	0.933	0.778	0.848	18
Macro avg	0.902	0.890	0.891	131
Weighted avg	0.906	0.893	0.895	131


The most dramatic improvement vs. Phase 1 is **surprised recall: 0.333 → 0.778** (more than doubled). Every class is now above 0.84 F1, whereas in Phase 1 *surprised* was at 0.462.

6.3 Confusion Matrix — Ensemble




Of 131 test samples, 117 are classified correctly (89.3%). The single largest off-diagonal block is **3 surprised** → **happy** (likely because surprised onsets share a rising-pitch profile with cheerful exclamations) and **4 neutral** → **happy** (a known difficulty: subdued enthusiasm vs. flat narration). All other off-diagonals are ≤ 2 samples.

6.4 Normalised Confusion Heatmap



Diagonal recalls (normalised): **angry 0.91, happy 0.96, neutral 0.86, sad 0.94, surprised 0.78**. Compare to Phase 1 where *surprised* sat at 0.33 — the addition of Δ -MFCC and Mel-spectrogram features is the decisive ingredient that lifted this class.

6.5 Random-Forest-Only Comparison



The Phase-2 RF alone reaches 81.7% — substantially better than Phase 1 RF (62.6%) thanks to the 135-dim feature set, but ~8 pp below the ensemble. The largest gap is on *happy* (RF recall 0.857 vs. ensemble 0.964), where MLP's non-linear decision surface complements RF's tree splits.

7. Error Analysis & Discussion

7.1 Remaining Sources of Error (14 mis-classifications out of 131)

Confusion pair (true → predicted)	Count	Diagnostic
neutral → happy	4	Subdued enthusiasm; overlapping mid-arousal pitch.
surprised → happy	3	Both share rising-pitch onsets; valence ambiguity.
angry → happy	2	Both high-arousal; differ in voice quality (creaky vs. modal).
sad → surprised	1	Likely a long-vowel utterance with marginal pitch rise.
surprised → sad	1	Speaker spoke softly, lowering energy below the surprised prototype.
Other (singletons)	3	No systematic pattern.

Diagnosis: the remaining errors concentrate on the *happy* column — the ensemble has slightly over-permissive *happy* boundaries. This is the natural follow-up target for Phase 3.

7.2 Why Δ-MFCC and Mel-spectrogram helped *Surprised* the most

Surprised utterances are characterised by a **fast pitch rise followed by a sudden fall** — a *dynamic* pattern, not a static prosodic level. Phase 1’s frame-aggregate statistics (mean F0, mean STE) blurred this dynamics into the same numerical range as *angry*. Δ-MFCC explicitly encodes frame-to-frame change of the spectral envelope, and the pooled Mel-band statistics retain enough perceptual frequency detail to distinguish the brief surprise burst from sustained angry shouting. The feature-importance ranking (§5.1) puts Δ-MFCC coefficients in five of the top six positions, quantifying this contribution.

7.3 Phase 3 Improvement Plan

Concrete strategies to push from 89.3% toward ≥ 92%:

- **Δ^2 -MFCC (acceleration coefficients)**: the second derivative captures *transitions of transitions* — useful for *neutral* ↔ *happy* whose first-derivative profiles overlap.
- **Voiced-frame ratio** and **speaking-rate** features: low-arousal classes (*sad*, *neutral*) have higher silence ratios — currently invisible to our static cepstral pipeline.
- **Class-specific decision thresholds**: after probability calibration (`CalibratedClassifierCV`), lower the *surprised* / *neutral* thresholds to bias the ensemble against the dominant *happy* attractor.
- **Data augmentation** for hard classes: pitch-shift ± 2 semitones and time-stretch $\pm 10\%$ the *surprised* and *neutral* training samples to triple their effective count.
- **Hierarchical classification**: a two-stage system that first decides arousal (high vs. low) and then valence within each branch could resolve the remaining *neutral* ↔ *happy* and *surprised* ↔ *happy* confusions, exploiting the known 2D structure of emotion space.

8. GitHub Link

<https://github.com/mustafa-akgul/Speech-Emotion-Classification>

The Phase 2 implementation is committed alongside the Phase 1 code so that the full evolution is reproducible from a single repository.

9. Resources

#	Resource	Usage
1	Librosa (https://librosa.org)	<code>feature.mfcc</code> , <code>feature.delta</code> , <code>feature.chroma_stft</code> , <code>feature.melspectrogram</code> , <code>feature.tonnetz</code> , <code>feature.spectral_flatness</code> , <code>effects.preemphasis</code> , <code>effects.trim</code>
2	scikit-learn (https://scikit-learn.org)	<code>RandomForestClassifier</code> , <code>MLPClassifier</code> , <code>VotingClassifier</code> , <code>RandomizedSearchCV</code> , <code>StratifiedKFold</code> , <code>StandardScaler</code> , <code>metrics</code>
3	XGBoost (https://xgboost.readthedocs.io)	<code>XGBClassifier</code> with sklearn-compatible API and <code>predict_proba</code> for ensemble

#	Resource	Usage
4	Schuller, B. et al. (2018) Comm. ACM	Feature taxonomy and ensemble motivation
5	Furui, S. (1986) ICASSP	Delta-cepstra rationale
6	Bou-Ghazale & Hansen (2000) IEEE T-SAP	Δ -MFCC under emotional stress
7	Eyben, F. et al. (2016) GeMAPS, IEEE T-AC	Feature-set design (flatness, harmonic, MFCC dynamics)
8	Livingstone & Russo (2018) RAVDESS, PLOS ONE	Baseline accuracy reference
9	Wang, K. et al. (2015) Multimed. Tools Appl.	Soft-voting ensemble in SER
10	Valerio Velardo — <i>The Sound of AI</i> YouTube series	Conceptual reference for MFCC, mel-spectrogram, chroma
11	Claude AI (claude.ai)	Code review and report drafting assistance

10. Prompts (AI Tool Usage)

#	Tool	Prompt summary	Usage context
1	Claude AI	"Phase 1 baseline 51-dim RF üzerine genişletilmiş feature seti (Δ -MFCC, Chroma, Mel-Spec, Tonnetz) ve üç-model soft-voting ensemble pipeline'ı yaz."	Phase 2 code skeleton generation
2	Claude AI	"RandomizedSearchCV ile RF, XGBoost ve MLP için ortak bir hiperparametre arama framework'ü kur."	Hyper-parameter optimisation design
3	Claude AI	"results.csv'den sklearn.classification_report ile per-class precision/recall/F1 hesapla ve tabloya dök."	Per-class metric computation
4	Claude AI	"Phase 1 vs Phase 2 model accuracy bar grafiği üret; her bar üstüne yüzde değeri yaz."	Comparison visualisation

#	Tool	Prompt summary	Usage context
5	Claude AI	"Phase 1 raporu şablonunu Phase 2'ye uyarla — Literature Review bölümü ekle, methodology Phase 2 değişikliklerini açıklasın."	Technical report drafting

11. Team Member Contributions

Ahmet Akın designed and implemented the extended feature pipeline. He researched the Δ -MFCC, Chroma, Tonnetz, Mel-spectrogram and Spectral-Flatness extractors, integrated them with the Phase 1 preprocessing chain (preemphasis + trim), and validated that the resulting 135-dim feature matrix is free of NaN/Inf values on the full dataset.

Berivan Demir built the model layer. She implemented the three base classifiers (RF, XGBoost, MLP), wrote the `RandomizedSearchCV` tuning routines, designed the soft-voting ensemble (including the `XGBWrapper` that bridges XGBoost's integer labels with sklearn's string-label `VotingClassifier`), and produced the confusion-matrix and feature-importance figures.

Mustafa Talha Akgül led error analysis and reporting. He computed the per-class precision/recall/F1 table from the ensemble predictions, produced the Phase 1 vs Phase 2 comparison chart, drafted the technical report (including the literature review and Phase 3 improvement plan), maintained the GitHub repository, and is responsible for the Phase 2 leaderboard submission.